

Section 4: Assembler & Toolchain Unified Plan (v1–v5)

Computer Ecosystem Project

January 29, 2026

Contents

1 Version Comparison (v1–v5)	3
2 Detailed Execution Plans	4
3 Plan of Execution (v1)	4
3.1 Purpose and Scope	4
3.2 Deliverables	4
3.3 CLI Contracts	4
3.4 Assembly Language Spec (Minimal)	4
3.5 Encoding Contract	4
3.6 Execution Checklist	5
3.7 Diagnostics Requirements	5
3.8 Definition of Done	5
4 Plan of Execution (v2)	6
4.1 Purpose and Scope	6
4.2 Deliverables	6
4.3 CLI Additions	6
4.4 Language Additions	6
4.5 Execution Checklist	6
4.6 New Diagnostics	6
4.7 Definition of Done	7
5 Plan of Execution (v3)	8
5.1 Purpose and Scope	8
5.2 Deliverables	8
5.3 Binary Layout Convention	8
5.4 Directive Semantics	8
5.5 Execution Checklist	8
5.6 New Diagnostics	9
5.7 Definition of Done	9
6 Plan of Execution (v4)	10
6.1 Purpose and Scope	10
6.2 Deliverables	10
6.3 BO Operand Syntax	10
6.4 Relocations (Minimal Set)	10
6.5 Execution Checklist	10

6.6	New Diagnostics	10
6.7	Definition of Done	11
7	Plan of Execution (v5)	12
7.1	Purpose and Scope	12
7.2	Deliverables	12
7.3	Two-Image Build Model	12
7.4	Execution Checklist	12
7.5	New Diagnostics	12
7.6	Definition of Done	13

1 Version Comparison (v1–v5)

Ver.	Primary Goal	Key Additions	Build Output
v1	Minimum end-to-end loop: .asm → .bin → run	Two-pass assembler, fixed 8-byte encoding, strict diagnostics, flat binary, minimal loader	Single flat .bin (base 0x0000)
v2	Enable structured programs	Stack + calls: PUSH8, POP8, CALL_ABS, RET, CMP; accept SP, FP; (opt) .org, .sym, disasm	Single flat .bin + (opt) .sym
v3	OS-shaped layout and control transfers	Vectors at 0x0000..0x00FF, reset at 0x0100; SYSCALL, IRET; directives .vectors, .vector, .entry, .text; image layout mode	Memory image .bin (vectors + text)
v4	Compiler-friendly milestone	BO addressing + wide ops; data directives + .align; sections .text, .data; minimal object format .o; minimal linker ld + relocations	.o objects + linked .bin + .map
v5	Protection-aware system build	Indirect control flow (JMP_R, JZ_R, JNZ_R, CALL_R); SETI, SETK, TRAP; dual-image builds (kernel and user); linker range checks + lints; trap-aware trace and debug	kernel.bin + user.bin + maps/metadata

2 Detailed Execution Plans

The following sections are included via \input. All files are intended to compile together as a single document.

3 Plan of Execution (v1)

3.1 Purpose and Scope

Goal. Deliver the minimum end-to-end loop:

```
.asm → assembler → flat .bin → loader → emulator run
```

Target: CPU v1 contract (fixed-size 8-byte instructions, strict validation, flat memory).

Out of Scope. Linker/object files, macros, includes, multi-file builds, sections, vectors, syscalls.

3.2 Deliverables

1. **Assembler (asm):** two-pass assembler producing a flat binary.
2. **Loader (loadbin):** loads .bin at base address.
3. **Integration demo (demo.asm)** that assembles and runs.
4. **Test suite:** unit + golden + end-to-end.
5. **Documentation:** syntax, mnemonics, errors, CLI.

3.3 CLI Contracts

```
# Assemble:  
asm input.asm -o program.bin  
  
# Optional:  
# --base 0x0000  
  
# Load into emulator memory:  
loadbin program.bin --base 0x0000
```

3.4 Assembly Language Spec (Minimal)

- Comments: ; to end of line
- Labels: name:
- Registers: R0..R15
- Integers: decimal and 0x hex

Supported mnemonics (v1): MOV_RI, MOV_RR, ADD, SUB, LOAD8_ABS, STORE8_ABS, JMP_ABS, JMP_REL, JZ_ABS, JZ_REL, HALT.

3.5 Encoding Contract

Each instruction is 8 bytes:

$$\text{opcode}(8) \mid \text{rd}(8) \mid \text{ra}(8) \mid \text{rb}(8) \mid \text{imm32}(32)$$

Little-endian for imm32.

3.6 Execution Checklist

1. Define instruction table: mnemonic → opcode + operand schema + zero-field constraints.
2. Implement lexer (tokens: ident, reg, number, punctuation, comments).
3. Implement parser (line → optional label + optional instruction).
4. Pass 1: assign addresses, build symbol table (PC starts at base, +8 per instruction).
5. Pass 2: resolve labels, compute REL offsets, validate ranges, emit bytes.
6. Write `.bin` (no header); implement `loadbin`.
7. Build `examples/v1/demo.asm`; add golden bytes test + emulator integration test.

3.7 Diagnostics Requirements

All errors must include file, line/column, category code, message, and (optional) hint.

3.8 Definition of Done

1. Demo assembles to stable bytes and runs to HALT.
2. Unit tests cover lexer/parser/encoder and key failure modes.
3. Docs describe syntax and CLI.

4 Plan of Execution (v2)

4.1 Purpose and Scope

Goal. Extend v1 to enable structured programs:

- Stack/calls: PUSH8, POP8, CALL_ABS, RET
- Comparison: CMP
- Assembly register names: SP, FP

Remain backwards compatible with v1 sources.

Out of Scope. Interrupt vectors, syscalls, MMIO layout (v3); sections and data (v4); user and kernel protection (v5).

4.2 Deliverables

1. Updated assembler with v2 mnemonics + SP, FP.
2. v2 example programs: `call_ret.asm`, `stack.asm`.
3. Updated tests (encoding + new diagnostics).
4. Recommended: `.sym` output; minimal disassembler.

4.3 CLI Additions

```
asm input.asm -o program.bin
# recommended:
# --sym program.sym
# --dump
```

4.4 Language Additions

New mnemonics: CMP, PUSH8, POP8, CALL_ABS, RET.

New register tokens: SP, FP (internally encoded using special selectors where ISA permits).

Optional directive: `.org <address>` (may require gap-fill with zeros in flat output).

4.5 Execution Checklist

1. Extend instruction table with v2 mnemonics and constraints (e.g., RET fields must be zero).
2. Update lexer/parser to accept SP, FP.
3. Implement (optional) `.org` in parser + pass1/pass2 with output gap-fill.
4. Implement encoder support for new instructions and strict field validation.
5. Add example programs and golden tests; run emulator integration (SP changes, CALL and RET returns).
6. (Recommended) Write `.sym`; (Recommended) implement disasm reading 8-byte chunks.

4.6 New Diagnostics

- **E_BAD_SPECIAL_REG:** SP, FP used where not allowed.
- **E_RET_NONZERO:** RET has nonzero fields.
- **E_ORG_BACKWARDS:** `.org` decreases PC (if forbidden).

4.7 Definition of Done

1. All v1 programs still assemble and run.
2. v2 call/stack examples assemble and run.
3. Tests validate new encodings and errors; optional sym/disasm works if implemented.

5 Plan of Execution (v3)

5.1 Purpose and Scope

Goal. Adopt OS-shaped layout and control transfers:

- Vector table at 0x0000..0x00FF
- Reset entry at PC=0x0100
- New mnemonics: `SYSCALL`, `IRET`
- MMIO conventions (console and timer)

Out of Scope. BO addressing + data sections + linker (v4); protection rules (v5).

5.2 Deliverables

1. Assembler supports `SYSCALL` and `IRET`.
2. Layout directives: `.vectors`, `.vector`, `.entry`, `.text`.
3. Image-mode output that places vectors and code at the proper addresses.
4. Loader defaults appropriate for reset at 0x0100.
5. Examples: vectors + syscall demo (+ timer IRQ demo if emulator supports).

5.3 Binary Layout Convention

Two modes:

- **Raw mode** (compat): behaves like v2 flat blob.
- **Image mode** (recommended): emits a memory image starting at 0x0000 with:
 - vectors at 0x0000
 - text at 0x0100
 - zero-filled gaps

5.4 Directive Semantics

```
.vectors
.vector 0, syscall_handler % 0x00: syscall vector
.vector 1, timer_handler % 0x01: timer interrupt vector

.entry start % reset begins at PC=0x0100
.text
```

Vector table contract (matches CPU v3). The vector table occupies 0x0000..0x00FF. Each vector entry is a 2-byte little-endian address (u16) stored at 0x0000 + 2**id*. Vector 0x00 is the syscall handler; vector 0x01 is the timer interrupt handler. On reset the CPU begins at PC=0x0100 (no reset vector).

5.5 Execution Checklist

1. Extend instruction table: `SYSCALL` (no operands; `rd=ra=rb=imm32=0`; syscall number in R0), `IRET` (all non-opcode fields zero).
2. Parse new directives and track state (inside `.vectors`, current section).
3. Implement image-mode layout engine: reserve 0x0000..0x00FF for vectors; emit vector entries at 0x0000; place `.text` at 0x0100; gap-fill zeros.
4. Pass 1: compute symbol addresses with layout rules.

5. Pass 2: emit bytes at correct offsets; patch vector slots as `u16` little-endian handler addresses at `0x0000 + 2*id` (valid ids `0x00..0x7F`).
6. Update loader/run defaults: load at `0x0000`, start execution at `0x0100` (reset entry).
7. Add demos and integration tests: SYSCALL enters handler, IRET returns.

5.6 New Diagnostics

- `E_BAD_VECTOR_ID`: vector id out of range.
- `E_VECTOR_NO_LABEL`: undefined handler label.
- `E_ENTRY_NOT_0100`: `.entry` label does not resolve to `0x0100` in v3 image mode.
- `E_SYSCALL_NONZERO`: SYSCALL has any nonzero non-opcode fields.
- `E_VECTOR_UNALIGNED`: handler address is not 8-byte aligned (instruction alignment).
- `E_OVERLAP_VECTORS`: code overlaps `0x0000..0x00FF` when forbidden.
- `E_IRET_NONZERO`: IRET has nonzero fields.

5.7 Definition of Done

1. v1 and v2 programs still assemble and run.
2. v3 image binaries place vectors/text correctly.
3. SYSCALL and IRET demo works (vector 0 dispatch + return) and tests cover directives + vector patching.

6 Plan of Execution (v4)

6.1 Purpose and Scope

Goal. Make the toolchain compiler-friendly:

- CPU v4 support: BO addressing + wide load and store + bitwise and shift
- Data directives and `.align`
- Minimal `.text`, `.data` sections
- Minimal object output (`.o`) + minimal linker (`ld`) + relocations

Out of Scope. Protection enforcement and user and kernel split rules (v5).

6.2 Deliverables

1. Assembler supports BO operands and v4 mnemonics.
2. Data directives: `.byte`, `.word16`, `.word32`, `.word64`, `.ascii`, `.asciz`, plus `.align`.
3. Sections: `.text` and `.data`.
4. Object output: `asm -c file.asm -o file.o`.
5. Linker: `ld a.o b.o -o program.bin` with relocations + `.map`.

6.3 BO Operand Syntax

Recommended:

```
LOAD32_BO R1, [R3 + 16]  
STORE8_BO [R4 + -1], R2
```

6.4 Relocations (Minimal Set)

- `R_ABS16`: write 16-bit absolute (LE)
- `R_ABS32`: write 32-bit absolute (LE)
- `R_REL32`: write 32-bit PC-relative (optional, if needed)

6.5 Execution Checklist

1. Add v4 ISA mnemonics and BO operand parsing; encode correctly; unit tests.
2. Add data directives emission (bytes + LE words + strings) and tests.
3. Add `.align N` padding with validation and tests.
4. Implement minimal sections in binary mode (two buffers: text and data; deterministic placement).
5. Add object output `-c`: section bytes + symbols + relocations; add `.global`, `.extern`.
6. Build minimal linker: merge sections, resolve globals, assign addresses, apply relocations, emit `.bin` + `.map`.
7. Add multi-file example: `lib.o` + `main.o` links and runs.

6.6 New Diagnostics

- `E_BAD_BO_ADDR`: malformed [Reg + Off].
- `E_ALIGN_ZERO`: invalid `.align 0`.
- `E_BAD_STRING`: invalid string escape.
- `E_DUP_GLOBAL`: duplicate global symbol at link time.
- `E_UNDEF_GLOBAL`: undefined external symbol at link time.
- `E_RELOC_RANGE`: relocation overflow.

6.7 Definition of Done

1. v1–v3 programs still build and run.
2. v4 BO + wide ops + data directives work with tests.
3. `asm -c` and `ld` link multi-file programs and run in emulator.

7 Plan of Execution (v5)

7.1 Purpose and Scope

Goal. Finish Section 4 roadmap alignment with protection-aware builds:

- Indirect control flow: `JMP_R`, `JZ_R`, `JNZ_R`, `CALL_R`
- Mode/trap ops: `SETI`, `SETK`, `TRAP`
- Dual-image build model: **kernel + user**
- Linker range enforcement + user MMIO lints
- Trap-aware trace and debug improvements (recommended)

Note. `JNZ_R` is introduced in v5 as the register-indirect complement to `JZ_R`.

7.2 Deliverables

1. Assembler supports v5 mnemonics and strict field validation.
2. Linker supports `-target=kernel` and `-target=user`.
3. Range checks:
 - user code and data must fit `-user-start..-user-end`
 - user output cannot reference MMIO absolute region (lint)
4. Minimal syscall and trap stubs library for userland (optional but strongly recommended).
5. End-to-end demos: kernel boots, transitions to user, user performs syscall or trap and returns.

7.3 Two-Image Build Model

- **Kernel image:**
 - owns vectors at `0x0000`
 - reset entry at `0x0100`
 - provides syscall and trap handlers
- **User image:**
 - linked into user-safe region
 - calls services via `SYSCALL` and/or `TRAP`
 - should not directly use MMIO absolute addresses

7.4 Execution Checklist

1. Add v5 mnemonics to instruction table and encoder (reg-only indirect jumps and calls; strict imm and zero-field rules).
2. Extend linker with `-target=kernel` and `-target=user` and separate layouts.
3. Implement region enforcement and MMIO lints for user builds; add negative tests.
4. Provide minimal user syscall stubs (`libsys`) + kernel handlers; buildable as `.o`.
5. Improve trace and debug: trap cause, EPC, BADADDR, mode transitions; symbolized breakpoints (recommended).
6. Create v5 end-to-end demo: kernel + user images, syscall prints, return path validated.

7.5 New Diagnostics

- **E_BAD_INDIRECT:** malformed indirect jump and call operands.
- **E_IMM_NONZERO:** instruction requires `imm=0` but found nonzero.
- **E_USER_MMIO:** user binary contains forbidden MMIO absolute access.

- **E_RANGE_OVERFLOW**: output exceeds allowed region.
- **E_BAD_TARGET**: user symbol resolves into kernel-only region.

7.6 Definition of Done

1. v1–v4 builds remain functional.
2. v5 instructions assemble/link/run.
3. Kernel and user images build with enforced ranges and lints.
4. End-to-end demo shows syscall or trap path and (optional) indirect calls.