

Section 5: Kernel Stub & Early Boot

Versioned Specifications (v1–v5)

Computer Ecosystem Project

January 29, 2026

Contents

1 Overview	3
2 Version Comparison Table (v1–v5)	4
2.1 High-level feature comparison	4
2.2 Compatibility summary (what stays invariant)	4
3 Specifications by Version	6
4 Section 5 v1: Kernel Stub & Early Boot (Minimal Boot Bridge)	6
4.1 Compatibility Contract (from earlier sections)	6
4.2 Reset and entry assumptions (v1)	6
4.3 Scope	6
4.4 Memory Plan (v1 recommended)	6
4.5 Control Flow (normative)	7
4.6 Kernel entry contract (v1 ABI)	7
4.7 Plan of Execution (v1)	7
4.8 Definition of Done (v1)	7
5 Section 5 v2: Kernel Stub & Early Boot (Deterministic Diagnostics)	8
5.1 Compatibility Contract	8
5.2 Scope	8
5.3 Boot stages (frozen IDs)	8
5.4 Panic contract (normative)	8
5.5 Panic codes (normative)	8
5.6 Sanity checks (mandatory)	9
5.7 Control flow (normative)	9
5.8 Definition of Done (v2)	9
6 Section 5 v3: Kernel Stub & Early Boot (Vectors + Reset @ 0x0100)	10
6.1 Compatibility Contract	10
6.2 Scope	10
6.3 Memory and image layout (normative)	10
6.4 Required vectors (v3)	10
6.5 Minimal syscall ABI (v3)	10
6.6 Interrupt model (v3/v4)	10
6.7 Boot stages and diagnostics continuity (normative)	11
6.8 Boot flow (normative)	11
6.9 v3 panic signature (normative)	11
6.10 Definition of Done (v3)	11

7	Section 5 v4: Kernel Stub & Early Boot (Boot ABI + BootInfo)	12
7.1	Compatibility Contract	12
7.2	Boot ABI (normative)	12
7.3	BootInfo placement (v4 new rule)	12
7.4	BootInfo v4 layout (normative)	12
7.5	Alignment and v4 memory ops (normative)	13
7.6	Syscall ABI refinement (v4)	13
7.7	Definition of Done (v4)	13
8	Section 5 v5: Kernel Stub & Early Boot (Protection + Dual Images)	14
8.1	Compatibility Contract	14
8.2	User-mode memory access rules (normative)	14
8.3	Reset convention (v5 alignment with CPU roadmap)	14
8.4	Recommended placement convention (non-normative)	14
8.5	Required vectors (v5)	14
8.6	BootInfo v5 extension (normative)	15
8.7	Kernel→user transition (normative)	15
8.8	Protection fault signature (normative)	15
8.9	Definition of Done (v5)	15

1 Overview

This document aggregates Section 5 (Kernel Stub & Early Boot) specifications across versions **v1–v5**. It includes:

- a comparison table showing what each version adds; and
- the full specification text for each version using \input.

2 Version Comparison Table (v1–v5)

2.1 High-level feature comparison

Capability / Rule	v1	v2	v3	v4	v5
Target CPU baseline	CPU v2	CPU v2	CPU v3	CPU v4	CPU v5
Output format	Flat <code>.bin</code>	Flat <code>.bin</code>	Image-mode (vectors + entry)	Image-mode + <code>.data</code> /fixed BootInfo	Dual-image (<code>kernel.bin</code> + <code>user.bin</code>)
Reset entry PC	0x0000	0x0000	0x0100	0x0100	0x0100
Vector table (0x0000..0x00FF)	— (not used)	— (not used)	Required	Required	Required + protection vector
Kernel stub responsibilities	Stack init + jump	+ sanity checks + panic convention	+ syscall/- timer entry paths	+ Boot ABI + BootInfo + v4 alignment discipline	+ kernel → user transition + protection boundary
Deterministic panic signature	Optional (minimal)	Required (R0..R4)	Required (vector- aware)	Required (BootInfo- aware)	Required (adds prot-fault signature)
Syscalls	No	No	Yes (<code>SYSCALL</code> via vector 0x00)	Yes (ABI stabilized + preservation rules)	Yes (user uses syscalls; kernel enforces boundary)
Timer interrupt	No	No	Yes (vector 0x01 minimal handler)	Yes	Yes
Boot ABI into <code>kernel_main</code>	Implicit	Implicit	Minimal (stack + jump)	Frozen: <code>R0=bootinfo_ptr</code>	Frozen + extended BootInfo for user image
BootInfo block	No	No	Optional (debug-only)	Required (v4 layout)	Required (v5 extended layout)
Protection / user mode	No	No	No	No	Yes (K flag + traps + dual-image)

2.2 Compatibility summary (what stays invariant)

- Instruction encoding remains fixed (8 bytes) across versions.
- v3+ reserves vectors at 0x0000..0x00FF and resets at 0x0100.
- Stack remains below reserved MMIO start 0xFE00, with recommended kernel stack top 0xFDFF.
- In v5, user-mode may access exttt0x0100..0xFDFF; only vectors (exttt0x0000..0x00FF) and

MMIO/reserved (exttt0xFE00..0xFFFF) are privileged. exttt0xF800..0xFDFF may be used by the kernel stack by convention.

- Syscall entry vector remains 0x00 from v3 onward.

3 Specifications by Version

4 Section 5 v1: Kernel Stub & Early Boot (Minimal Boot Bridge)

Purpose

A **kernel stub** is the first OS-controlled code executed after CPU reset. In this project its role is intentionally narrow:

- establish the minimum execution environment that higher-level kernel code assumes;
- transfer control deterministically to the kernel entry function;
- remain forward-compatible with later CPU/toolchain versions.

Non-goal: the stub is not “the kernel”; it is the bridge that makes kernel code possible.

Version intent

v1 is the most minimalistic version that still works for this project: initialize a valid stack and jump to `kernel_main`. Everything else (diagnostics, vectors, syscalls, protection) is deferred to later versions.

4.1 Compatibility Contract (from earlier sections)

- **CPU baseline: v2.** v2 is the earliest version that makes SP/FP meaningful for a real stub.
- **Instruction encoding fixed at 8 bytes** across all versions (forward-compatibility invariant).
- **Address space:** 64KB RAM, byte-addressed, 0x0000..0xFFFF.
- **Little-endian** (relevant for call/return conventions later).
- **Assembler baseline: v2,** output is a **flat binary .bin** loaded at 0x0000.

4.2 Reset and entry assumptions (v1)

The CPU v2 reset convention provides a deterministic baseline:

- **PC** starts at 0x0000.
- **GPRs** start at 0; **Z** starts cleared.
- **SP** and **FP** start at 0xFDFF.

v1 still explicitly sets SP and FP to avoid depending on reset defaults.

4.3 Scope

In scope (required):

1. Initialize stack (SP) and frame pointer (FP).
2. Transfer control to `kernel_main`.

Out of scope: vectors, syscalls, interrupts, MMIO, allocators, protection/user mode.

4.4 Memory Plan (v1 recommended)

Region	Address range	Usage
Kernel image (flat)	0x0000 .. <end>	Stub + early kernel code
Free RAM	<end> .. 0xF7FF	Unused in v1
Kernel stack	0xF800 .. 0xFDFF	Downward-growing; top at 0xFDFF
Reserved future region	0xFE00 .. 0xFFFF	Keep unused (future MMIO)

4.5 Control Flow (normative)

```
_start (PC=0x0000):
    SP := 0xFDFF
    FP := 0xFDFF
    JMP_ABS kernel_main
```

4.6 Kernel entry contract (v1 ABI)

On entry to `kernel_main`:

- SP points to the stack-top (recommended 0xFDFF);
- FP = SP;
- no interrupts, syscalls, or IO are assumed.

4.7 Plan of Execution (v1)

Deliverables:

- `kstub_v1.asm` implementing stack init + jump;
- `kernel_main.asm` proving control transfer (deterministic HALT is enough);
- flat-binary build using assembler v2.

4.8 Definition of Done (v1)

v1 is complete when reset always reaches `kernel_main` and ends deterministically.

5 Section 5 v2: Kernel Stub & Early Boot (Deterministic Diagnostics)

Purpose

v2 preserves v1's minimal boot bridge (stack init + jump) and adds a frozen, deterministic early-boot diagnostics model:

If early boot fails, it fails loudly and repeatably via a panic signature.

5.1 Compatibility Contract

- **CPU baseline:** v2 (stack model exists).
- **Assembler baseline:** v2, output stays flat .bin loaded at 0x0000.
- No vectors, syscalls, or MMIO are used in v2.

5.2 Scope

Adds to v1:

- boot stage identifiers;
- minimal sanity checks;
- a frozen panic contract (register signature + panic codes).

5.3 Boot stages (frozen IDs)

Stage ID	Meaning
0x01	Entered <code>_start</code>
0x02	Stack initialized
0x03	Sanity checks running
0x04	Transfer to <code>kernel_main</code> imminent

5.4 Panic contract (normative)

On unrecoverable early-boot failure, the stub must set and then halt:

Register	Meaning
R0	Panic code
R1	Stage ID
R2	Location identifier (PC snapshot if available, else label constant)
R3	SP snapshot
R4	FP snapshot

5.5 Panic codes (normative)

Code	Meaning
0x10	SP out of allowed range
0x11	SP overlaps reserved future region (0xFE00..0xFFFF)
0x12	SP overlaps kernel code region (unsafe)
0x13	Alignment violation (PC/target not 8-byte aligned)
0x14	Kernel entry invalid/out of bounds

5.6 Sanity checks (mandatory)

Before jumping to `kernel_main`, v2 must verify:

1. SP is within `[STACK_BOTTOM..STACK_TOP]` (recommended `0xF800..0xFDFF`);
2. SP is below `0xFE00` (reserved for future MMIO);
3. `kernel_main` is valid and 8-byte aligned.

5.7 Control flow (normative)

```
_start:  
    stage := 0x01  
    SP := STACK_TOP  
    FP := STACK_TOP  
    stage := 0x02  
  
    stage := 0x03  
    run sanity checks  
    if fail -> PANIC(code, stage, location)  
  
    stage := 0x04  
    JMP_ABS kernel_main
```

5.8 Definition of Done (v2)

- normal boot reaches `kernel_main` exactly like v1;
- at least one deliberate misconfiguration triggers a correct panic signature.

6 Section 5 v3: Kernel Stub & Early Boot (Vectors + Reset @ 0x0100)

Purpose

v3 transitions the project into an OS-shaped boot model:

- vector table at 0x0000..0x00FF;
- reset entry at PC=0x0100;
- SYSCALL enters vector 0x00 and returns via IRET;
- timer interrupt enters vector 0x01 and returns via IRET.

6.1 Compatibility Contract

- **CPU baseline: v3.**
- **Toolchain baseline: v3 image-mode** capable of placing vectors and a fixed entry.
- Vector entries are **u16 little-endian addresses** at $0x0000 + 2*id$.
- Reset begins at 0x0100 (not 0x0000).

6.2 Scope

Adds to v2:

- required vector table and handlers;
- syscall entry/return proof-of-life;
- timer IRQ entry/return proof-of-life.

Still out of scope: protection/user mode and dual-image boot.

6.3 Memory and image layout (normative)

Region	Range	Rule
Vectors	0x0000..0x00FF	Present; kernel-owned
Reset entry	0x0100	<code>_start</code> placed here
Reserved MMIO	0xFE00..0xFFFF	Reserved for devices

6.4 Required vectors (v3)

ID	Handler	Requirement
0x00	<code>syscall_handler</code>	Must end with IRET
0x01	<code>timer_handler</code>	Must end with IRET

6.5 Minimal syscall ABI (v3)

- R0: syscall number / return value
- R1..R3: up to three arguments

Minimal proof-of-life syscalls:

- 1 = `sys_putc` (write R1 to MMIO console or debug RAM);
- 2 = `sys_halt`.

6.6 Interrupt model (v3/v4)

v3 introduces asynchronous interrupts (used by the timer IRQ and optionally other future IRQs).

There is no architectural interrupt-enable/mask bit in v3/v4. In other words, software cannot selectively disable interrupts at the ISA level until v5.

Normative entry/return rules:

- On IRQ entry, the CPU transfers control to the handler address read from the vector table (e.g., timer IRQ uses vector 0x01).
- The return address is saved in EPC (per the CPU roadmap), and the handler returns with IRET, restoring PC from EPC.
- Any additional trap metadata (e.g., CAUSE, BADADDR) is handled as defined by the CPU roadmap for the target version.

6.7 Boot stages and diagnostics continuity (normative)

For deterministic diagnostics, v3 **inherits the v2 boot stage IDs unchanged** for use in the panic signature. Implementations may report the same stage values as v2 (e.g., 0x01..0x04) when failing during early boot.

6.8 Boot flow (normative)

```
Reset:
PC := 0x0100

_start @ 0x0100:
SP := 0xFDFF
FP := 0xFDFF
JMP_ABS kernel_main
```

6.9 v3 panic signature (normative)

On unrecoverable fault, set and halt:

Reg	Meaning
R0	Panic code
R1	Stage ID
R2	Location identifier (EPC snapshot if available, else constant/label)
R3	SP snapshot
R4	FP snapshot

6.10 Definition of Done (v3)

- image boots at 0x0100;
- SYSCALL enters vector 0x00 and returns via IRET;
- timer IRQ enters vector 0x01 and returns via IRET.

7 Section 5 v4: Kernel Stub & Early Boot (Boot ABI + BootInfo)

Purpose

v4 keeps v3's layout and vectors unchanged, but stabilizes early boot as a platform for real kernel development by introducing:

- a frozen **Boot ABI** into `kernel_main`;
- a structured **BootInfo** block passed to the kernel;
- v4 alignment discipline enabling safe 16/32/64-bit memory operations and base+offset addressing.

7.1 Compatibility Contract

- **CPU baseline: v4.**
- v3 vector layout is preserved (vectors at 0x0000..0x00FF, reset 0x0100).
- Toolchain supports image-mode plus `.data` (or fixed placement) for `BootInfo`.

7.2 Boot ABI (normative)

Kernel entry is:

```
kernel_main(bootinfo_ptr)
```

Passing convention:

- R0 = `bootinfo_ptr`;
- SP initialized to kernel stack top (recommended 0xFDFF);
- FP = SP;

7.3 BootInfo placement (v4 new rule)

- **Recommended fixed address:** `BOOTINFO_ADDR` = 0x0200.
- **No-overlap rule (normative):** the kernel image (code/data/bss/stack init) must not overwrite the `BootInfo` block: reserve the range `[BOOTINFO_ADDR .. BOOTINFO_ADDR + size]` as kernel-owned metadata. If a fixed address is used, the linker/loader must place kernel sections so they do not overlap this range.
- Must be at least 8-byte aligned.
- Multi-byte fields are little-endian and naturally aligned.

7.4 BootInfo v4 layout (normative)

Offset	Type	Name	Meaning
0x00	u16	<code>magic</code>	0xB007
0x02	u8	<code>version</code>	4
0x03	u8	<code>size</code>	total size (recommended 0x20)
0x04	u16	<code>mem_size</code>	65536
0x06	u16	<code>flags</code>	feature flags
0x08	u16	<code>kernel_entry</code>	0x0100
0x0A	u16	<code>vector_base</code>	0x0000
0x0C	u16	<code>stack_top</code>	0xFDFF
0x0E	u16	<code>stack_bottom</code>	0xF800
0x10	u16	<code>reserved_mmio</code>	0xFE00
0x12	u16	<code>dbg_out</code>	optional debug RAM addr
0x14	u32	<code>boot_time</code>	optional (0 if unused)
0x18	u32	<code>reserved0</code>	reserved

0x1C	u32	reserved1	reserved
------	-----	-----------	----------

7.5 Alignment and v4 memory ops (normative)

- 16-bit ops: addr mod 2 = 0; 32-bit: mod 4 = 0; 64-bit: mod 8 = 0.
- v4 boot must include a minimal proof routine that:
 - reads BootInfo fields using base+offset addressing, and
 - clears a small aligned region using 64-bit stores.

7.6 Syscall ABI refinement (v4)

Syscall register convention remains: R0 syscall #, R1..R3 args, R0 return. v4 additionally freezes that syscall handlers must preserve callee-saved registers.

7.7 Definition of Done (v4)

- `kernel_main` receives R0=0x0200;
- BootInfo is validated (magic/version/size) and consumed via base+offset;
- a wide-store aligned memory clear succeeds;
- syscall/timer handlers still round-trip via IRET.

8 Section 5 v5: Kernel Stub & Early Boot (Protection + Dual Images)

Purpose

v5 adds the final foundational OS boundary:

kernel/user privilege separation and a dual-image boot model.

In v5:

- kernel boots privileged (K=1), owns vectors and BootInfo;
- user code runs unprivileged (K=0) in a user-safe region;
- user requests kernel services via **SYSCALL** (vector 0x00);
- illegal user access to vectors/MMIO triggers a protection trap to a dedicated vector.

8.1 Compatibility Contract

- **CPU baseline:** v5 (K flag + protection traps exist).
- v3/v4 vector layout preserved: vectors 0x0000..0x00FF, reset 0x0100.
- Toolchain v5 supports **dual-image outputs**: `kernel.bin` + `user.bin`.
- Boot ABI from v4 remains: `kernel_main(bootinfo_ptr)` with R0=bootinfo_ptr.

8.2 User-mode memory access rules (normative)

- **Vectors:** 0x0000..0x00FF (privileged-only)
- **User memory window:** 0x0100..0xFDFF (user-accessible when K=0)
- **Reserved/MMIO:** 0xFE00..0xFFFF (privileged-only)

When K=0, any fetch or data access into 0x0000..0x00FF or 0xFE00..0xFFFF triggers a protection trap to vector 0x03.

Kernel stack convention (non-normative). The kernel *may* use high RAM (e.g., 0xF800..0xFDFF) as its stack/critical scratch space, but this is a **software convention** rather than a hardware-enforced exclusion in v5.

8.3 Reset convention (v5 alignment with CPU roadmap)

On reset, the CPU state follows the v5 roadmap:

- PC=0x0100, SP=FP=0xFDFF, EPC=0
- FLAGS.ZNVC=0, FLAGS.I=1, FLAGS.K=1
- CAUSE=0, BADADDR=0

8.4 Recommended placement convention (non-normative)

To keep a clean software separation, the kernel *should* place its code/data in low memory (commonly below 0x8000) and load the user image at a higher base (recommended `user_base=0x8000`). This is a placement convention, not a CPU-enforced boundary; the kernel must therefore avoid relying on secrecy for data stored in the user-accessible RAM window.

8.5 Required vectors (v5)

ID	Handler	Requirement
0x00	<code>syscall_handler</code>	Must end with IRET
0x01	<code>timer_handler</code>	Must end with IRET
0x03	<code>prot_fault_handler</code>	Deterministic termination and fault signature

8.6 BootInfo v5 extension (normative)

BootInfo keeps the same header and extends with user-image fields (version=5, larger size). Key additional fields (conceptually):

- `user_base` (must be within 0x0100..0xFDFF; recommended 0x8000),
- `user_size`,
- `user_entry`,
- `user_stack_top` (recommended 0xFDFF),
- `prot_vector` (0x03).

8.7 Kernel→user transition (normative)

After kernel initialization, kernel must:

1. validate BootInfo v5 user fields (0x0100..0xFDFF range + 8-byte alignment);
2. prepare user stack and entry;
3. set `K := 0`;
4. jump to `user_entry` in the RAM window.

8.8 Protection fault signature (normative)

On protection fault (vector 0x03), kernel must record deterministically:

Reg	Meaning
R0	Fault code (0x50 protection violation)
R1	Vector ID (0x03)
R2	Location identifier (EPC snapshot if available, else constant)
R3	SP snapshot
R4	FP snapshot

8.9 Definition of Done (v5)

- dual-image boot reaches user entry in user mode (`K=0`);
- user syscall enters kernel via vector 0x00 and returns via `IRET`;
- user access to vectors/MMIO traps to vector 0x03 with correct fault signature.